

# AE 3: Duke Forest houses

## Checking model conditions

Add your name here

### Packages

```
library(tidyverse)
library(tidymodels)
library(openintro)
library(knitr)
```

### Predict sale price from area

```
df_fit <- linear_reg() %>%
  set_engine("lm") %>%
  fit(price ~ area, data = duke_forest)

tidy(df_fit) %>%
  kable(digits = 2)
```

| term        | estimate  | std.error | statistic | p.value |
|-------------|-----------|-----------|-----------|---------|
| (Intercept) | 116652.33 | 53302.46  | 2.19      | 0.03    |
| area        | 159.48    | 18.17     | 8.78      | 0.00    |

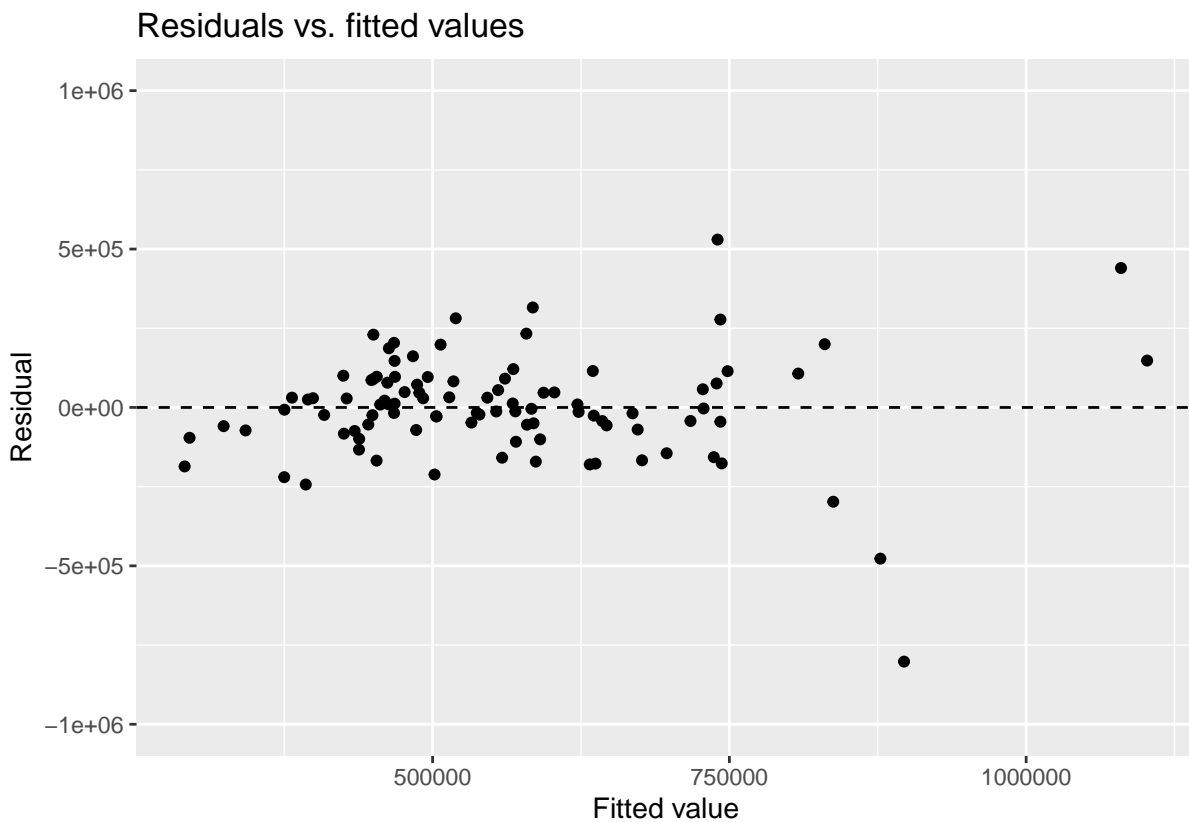
## Model conditions

### Exercise 1

The following code produces the residuals vs. fitted values plot for this model. Comment out the layer that defines the y-axis limits and re-create the plot. How does the plot change? Why might we want to define the limits explicitly?

```
df_aug <- augment(df_fit$fit)

ggplot(df_aug, aes(x = .fitted, y = .resid)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed") +
  ylim(-1000000, 1000000) +
  labs(
    x = "Fitted value", y = "Residual",
    title = "Residuals vs. fitted values"
  )
)
```



## Exercise 2

Improve how the values on the axes of the plot are displayed by modifying the code below.

```
ggplot(df_aug, aes(x = .fitted, y = .resid)) +  
  geom_point() +  
  geom_hline(yintercept = 0, linetype = "dashed") +  
  ylim(-1000000, 1000000) +  
  labs(  
    x = "Fitted value", y = "Residual",  
    title = "Residuals vs. fitted values"  
  )
```

